

# 第四章 数据流挖掘（上）

主讲：陈爱国

大数据分析 with 挖掘



# 主要内容

- 4.1 流数据模型
- 4.2 流数据抽样
- 4.6 窗口内的计数问题
- 4.3 流过滤
- 4.4 流中独立元素的数目统计
- 4.5 矩估计
- 4.7 衰减窗口

# New Topic: Infinite Data

## High dim. data

Locality  
sensitive  
hashing

Clustering

Dimensional  
ity  
reduction

## Graph data

PageRank,  
SimRank

Community  
Detection

Spam  
Detection

## Infinite data

Filtering  
data  
streams

Queries on  
streams

Web  
advertising

## Machine learning

SVM

Decision  
Trees

Perceptron,  
kNN

## Apps

Recommen  
der systems

Association  
Rules

Duplicate  
document  
detection

# So far

- So far we have worked datasets or data bases where all data is available
- In contrast, in **data streams**, data arrives one element at a time often at a **rapid rate** that:
  - If it is not **processed immediately** it is lost forever.
  - It is not feasible to **store** it all

# Data Streams

- In many data mining situations, we do not know the entire data set in advance
- **Stream Management** is important when the input rate is controlled **externally**:
  - Google queries
  - Twitter or Facebook status updates
- We can think of the **data** as **infinite** and **non-stationary** (无穷无尽, 且非平稳的, 数据分布会动态变化)

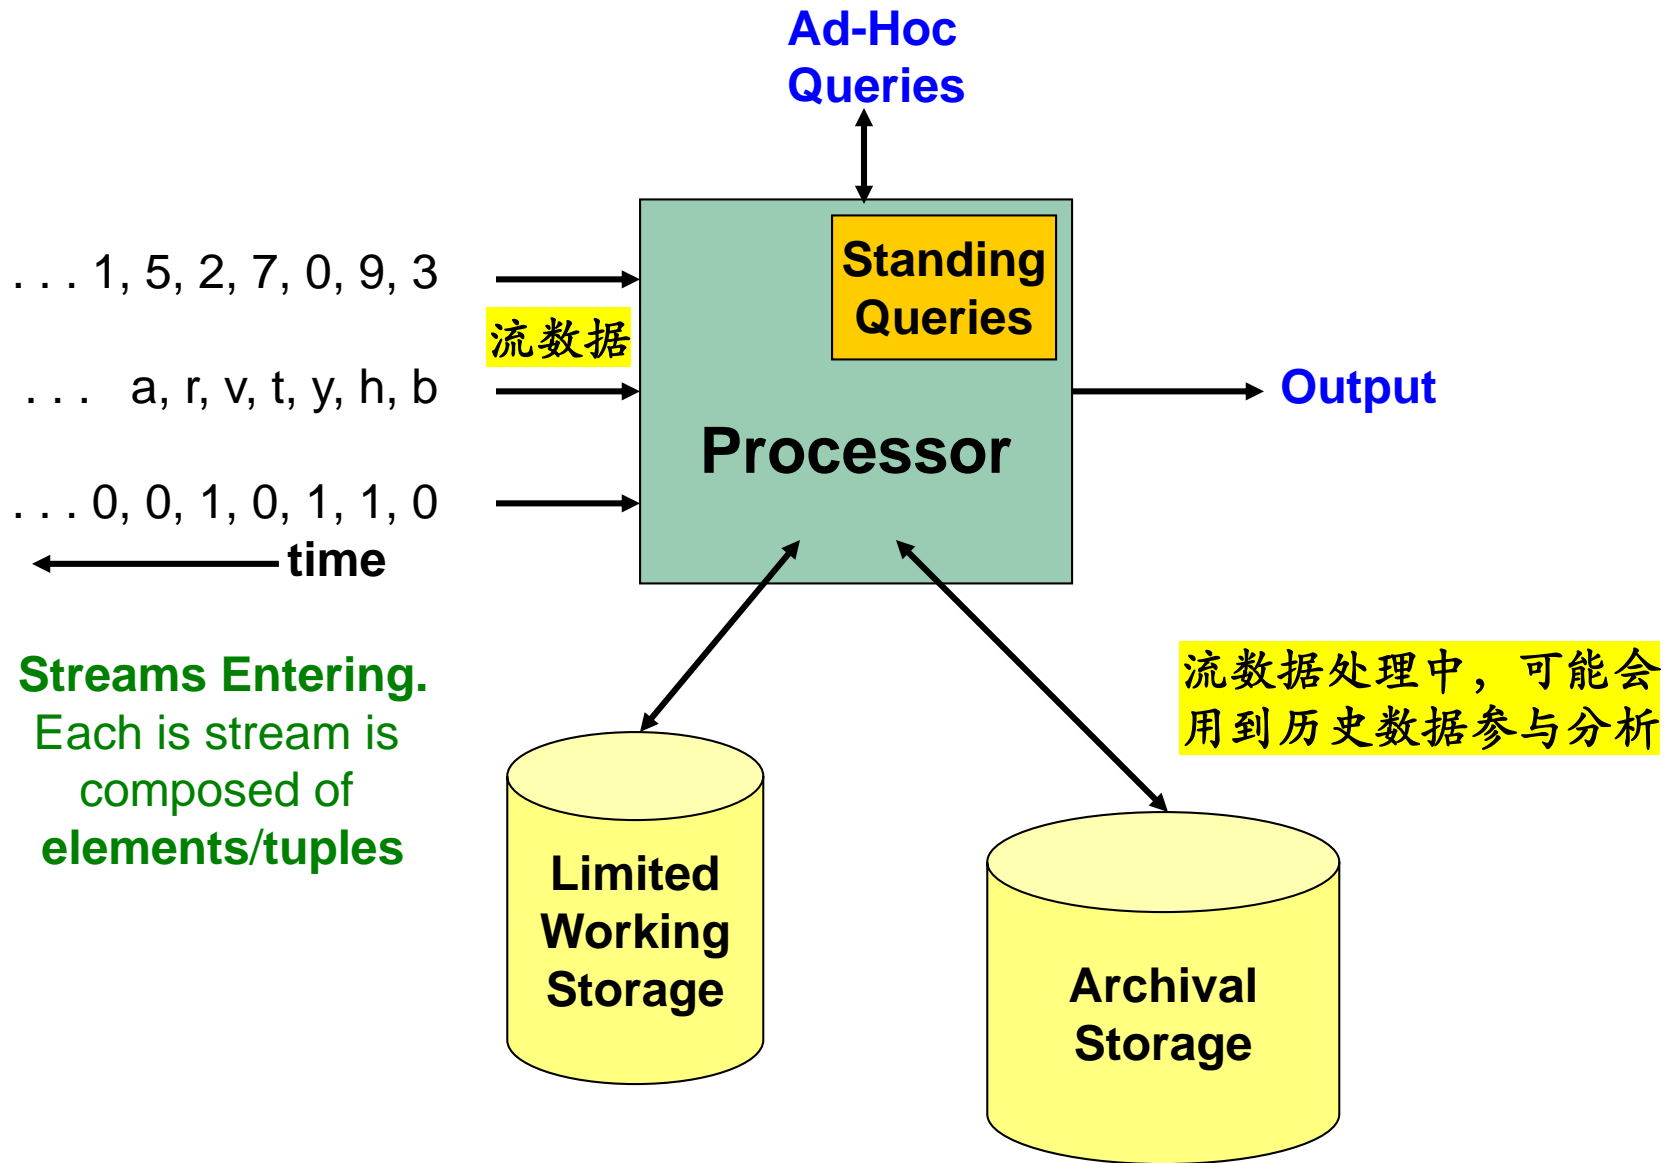
# The Stream Model

- Input **elements** enter at a rapid rate, at one or more input ports (i.e., **streams**)
  - **We call elements of the stream tuples**
- **The system cannot store the entire stream accessibly**
- **Q: How do you make critical calculations about the stream using a limited amount of (secondary) memory?**

# Side note: SGD is a Streaming Alg.

- **Stochastic Gradient Descent (SGD, 随机梯度下降)** is an example of a stream algorithm
- **In Machine Learning we call this: Online Learning**
  - Allows for modeling problems where we have a continuous stream of data
  - We want an algorithm to learn from it and slowly adapt to the changes in data
- **Idea: Do slow updates to the model**
  - **SGD** (SVM, Perceptron) makes small updates
  - **So:** First train the classifier on training data.
  - **Then:** For every example from the stream, we slightly update the model (using small learning rate)

# General Stream Processing Model





# Problems on Data Streams

- **Types of queries one wants on answer on a data stream:** (we'll do these today)
  - **Sampling data from a stream** 采样分析
    - Construct a random sample
  - **Queries over sliding windows** 滑动窗口查询
    - Number of items of type  $x$  in the last  $k$  elements of the stream

# Problems on Data Streams

- **Types of queries one wants on answer on a data stream:** (we'll do these next time)
  - **Filtering a data stream**
    - Select elements with property  $x$  from the stream
  - **Counting distinct elements**
    - Number of distinct elements in the last  $k$  elements of the stream
  - **Estimating moments**
    - Estimate avg./std. dev. of last  $k$  elements
  - **Finding frequent elements**

# Applications (1)

- **Mining query streams**

- Google wants to know what queries are more frequent today than yesterday

- **Mining click streams**

- Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour

- **Mining social network news feeds**

- E.g., look for trending topics on Twitter, Facebook

# Applications (2)

- **Sensor Networks**

- Many sensors feeding into a central controller

- **Telephone call records**

- Data feeds into customer bills as well as settlements between telephone companies

- **IP packets monitored at a switch**

- Gather information for optimal routing
- Detect denial-of-service attacks

# Sampling from a Data Stream: Sampling a fixed proportion

As the stream grows the sample  
also gets bigger

# Sampling from a Data Stream

- Why is this important?
  - Since **we can not store the entire stream**, a representative **sample** can act like the stream
- **Two different problems:**
  - (1) Sample a **fixed proportion** of elements in the stream (say 1 in 10) 固定比例采用, 如1/10
  - (2) Maintain a **random sample of fixed size** over a potentially infinite stream 固定规模的随机采样
    - At any “time”  $k$  we would like a random sample of  $s$  elements of the stream  $1\dots k$ 
      - What is the property of the sample we want to maintain?  
For all time steps  $k$ , each of  $k$  elements seen so far must have equal probability of being sampled

# Sampling a Fixed Proportion

- **Problem 1: Sampling a fixed proportion**
  - E.g. sample 10% of the stream
  - As stream gets bigger, sample gets bigger
- **Naïve solution:**
  - Generate a random integer in **[0...9]** for each query
  - Store the query if the integer is **0**, otherwise discard
- **Any problem with this approach?**
  - We have to be very careful what query we answer using this sample

需要结合业务需求场景，以体现模型价值和判断合理性

# Problem with Naïve Approach

- **Scenario:** Search engine query stream
  - **Stream of tuples:** (user, query, time)
  - **Question:** What fraction of unique queries by an average user are duplicates?
    - Suppose each user issues  $x$  queries once and  $d$  queries twice in one month
    - total of  $x+2d$  query instances
    - then the correct answer to the query is  $d/(x+d)$



# Problem with Naïve Approach

- **Scenario:** Search engine query stream
  - **Proposed solution:** We keep 10% of the queries
    - Sample will contain  $(x+2d)/10$  elements of the stream
    - Sample will contain  $d/100$  pairs of duplicates
      - $d/100 = 1/10 \cdot 1/10 \cdot d$
    - There are  $(10x+19d)/100$  unique elements in the sample
      - $(x+2d)/10 - d/100 = (10x+19d)/100$
  - So the sample-based answer is  $\frac{\frac{x}{10} + \frac{\frac{d}{100}}{100} + \frac{18d}{100}}{\frac{d}{100}} = \frac{d}{10x + 19d}$   
 $\neq \frac{d}{(x+d)}$

# Solution: Sample Users

## Solution:

- Pick  $1/10^{\text{th}}$  of **users** and take all their searches in the sample
- Use a hash function that hashes the user name or user id uniformly into 10 buckets

# Generalized Solution

- **Stream of tuples with keys:**
  - Key is some subset of each tuple's components
    - e.g., tuple is (user, search, time); key is **user**
  - Choice of key depends on application
- **To get a sample of  $a/b$  fraction of the stream:**
  - Hash each tuple's key uniformly into  **$b$**  buckets
  - Pick the tuple if its hash value is at most  **$a$**




Hash table with  **$b$**  buckets, pick the tuple if its hash value is at most  **$a$** .

**How to generate a 30% sample?**

Hash into  $b=10$  buckets, take the tuple if it hashes to one of the first 3 buckets

# Sampling from a Data Stream: Sampling a fixed-size sample

The sample is of fixed size



# Maintaining a fixed-size sample

- **Problem 2: Fixed-size sample**
- **Suppose we need to maintain a random sample  $S$  of size exactly  $s$  tuples** 仅保存 $s$ 个元组
  - E.g., main memory size constraint
- **固定大小抽样，希望达到如下效果：**
  - **Suppose at time  $n$  we have seen  $n$  items**
  - **Each item is in the sample  $S$  with equal prob.  $s/n$**

# Maintaining a fixed-size sample

例如：

**How to think about the problem: say  $s = 2$**

Stream: a x c y z k c d e g...

At  $n = 5$ , each of the first 5 tuples is included in the sample  $\mathbf{S}$  with equal prob.

At  $n = 7$ , each of the first 7 tuples is included in the sample  $\mathbf{S}$  with equal prob.

看似理想，但不现实的做法是：

存下  $n$  个元组（ $n$  是不断增大的），然后随机选择  $s$  个。

# Solution: Fixed Size Sample

- **Algorithm** (a.k.a. Reservoir Sampling, 水塘抽样)
  - Store all the first  $s$  elements of the stream to  $S$
  - Suppose we have seen  $n-1$  elements, and now the  $n^{th}$  element arrives ( $n > s$ )
    - With probability  $s/n$ , keep the  $n^{th}$  element, else discard it
    - If we picked the  $n^{th}$  element, then it replaces one of the  $s$  elements in the sample  $S$ , picked uniformly at random
- **Claim:** This algorithm maintains a sample  $S$  with the desired property:
  - After  $n$  elements, the sample contains each element seen so far with probability  $s/n$

# Proof: By Induction

- **We prove this by induction:**
  - Assume that after  $n$  elements, the sample contains each element seen so far with probability  $s/n$
  - We need to show that after seeing element  $n+1$  the sample maintains the property
    - Sample contains each element seen so far with probability  $s/(n+1)$  注意 $n$ 是动态变化的, 需要存储 $n$ 数值
- **Base case:**
  - After we see  $n=s$  elements the sample  $S$  has the desired property
    - Each out of  $n=s$  elements is in the sample with probability  $s/s = 1$



# Proof: By Induction

- **Inductive hypothesis:** After  $n$  elements, the sample  $S$  contains each element seen so far with prob.  $s/n$
- **Now element  $n+1$  arrives**
- **Inductive step:** For elements already in  $S$ , probability that the algorithm keeps it in  $S$  is:
$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$
- So, at time  $n$ , tuples in  $S$  were there with prob.  $s/n$
- Time  $n \rightarrow n+1$ , tuple stayed in  $S$  with prob.  $n/(n+1)$
- So prob. tuple is in  $S$  at time  $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

窗口内计数问题

Queries over a  
(long) Sliding Window

# Sliding Windows

- A useful model of stream processing is that queries are about a **window** of length  **$N$**  – the  **$N$**  most recent elements received
- **Interesting case:**  **$N$  is so large** that the data cannot be stored in memory, or even on disk
  - Or, there are so many streams that windows for all cannot be stored
- **Amazon example:**
  - For every product  **$X$**  we keep 0/1 stream of whether that product was sold in the  **$n$** -th transaction
  - We want answer queries, how many times have we sold  **$X$**  in the last  **$k$**  sales ( **$k$**  小于  **$n$** , 在窗口之内)

# Sliding Window: 1 Stream

## ■ Sliding window on a single stream:

N = 6

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past

Future →

# Counting Bits (1)

## ■ Problem:

- Given a stream of **0s** and **1s**
- Be prepared to answer queries of the form  
**How many 1s are in the last  $k$  bits?** where  $k \leq N$

*$N$ 是窗口大小,  $k$ 是任务中设定的一个参数, 可变*

## ■ Obvious solution:

Store the most recent  $N$  bits

- When new bit comes in, discard the  $N+1^{\text{st}}$  bit

0 1 0 0 1 1 0 1 1 1 0 1 0 1 1 0 1 1 0 1 1 0      Suppose  $N=6$

← Past                      Future →

# Counting Bits (2)


- You can not get an exact answer without storing the entire window
- **Real Problem:** 真正的问题是无法存下  $N$  bits 数据  
**What if we cannot afford to store  $N$  bits?**
  - E.g., we're processing 1 billion streams and  $N = 1$  billion



不能完整记录  $N$  个元组，同时希望得到准确结果

# An attempt: Simple solution

- Q: How many 1s are in the last  $N$  bits?



- **Simple solution**

- **Maintain 2 counters:**

- $S$ : number of 1s from the beginning of the stream
- $Z$ : number of 0s from the beginning of the stream

- **How many 1s are in the last  $N$  bits?**  $N \cdot \frac{S}{S+Z}$

- A simple solution that does not really solve our problem
- 它用了一致性假设(Uniformity assumption), 实际上数据流可能是非一致性的(non-uniform)
  - the distribution changes over time

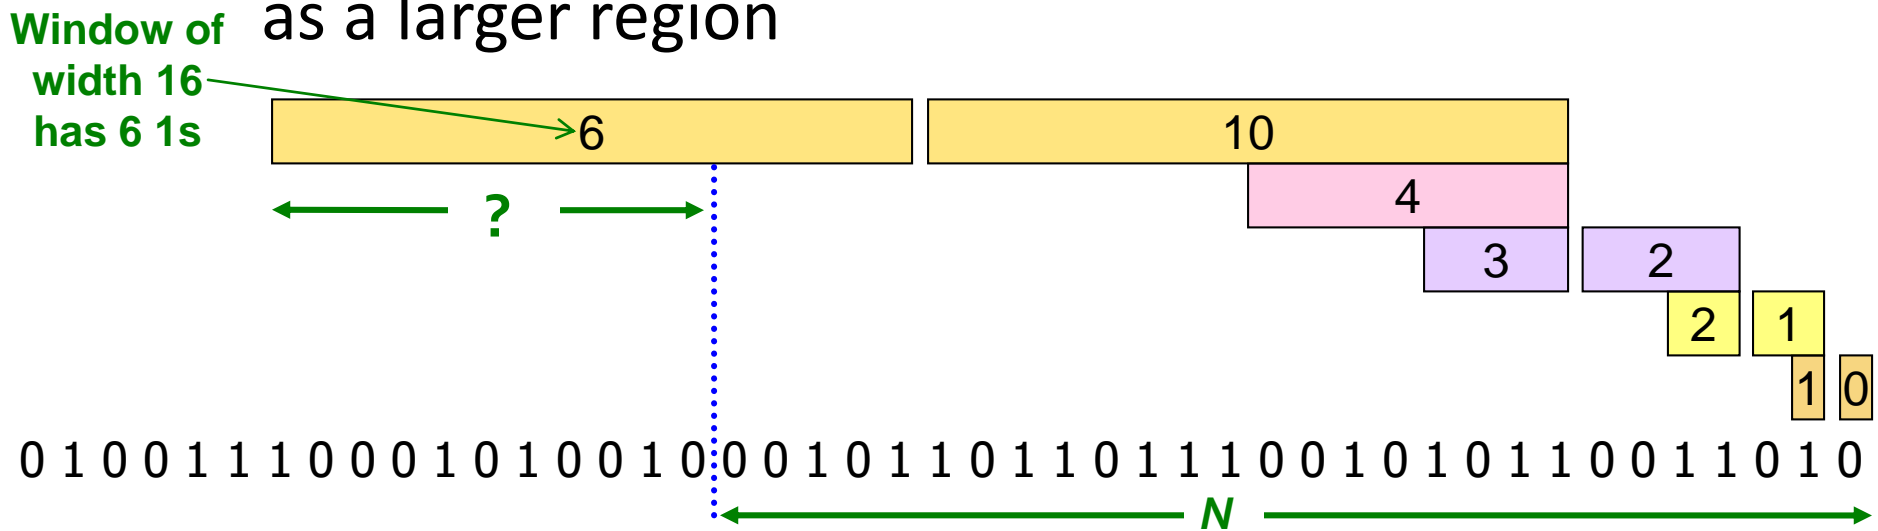
# DGIM Method

- Datar-Gionis-Indyk-Motwani Algorithm
  - DGIM solution that does not assume uniformity
- 用  $O(\log^2 N)$  位，表示大小为  $N$  的窗口
- 窗口内 1 的数量的估计误差，不超过 50%
- 后续的改进算法，可以不断降低错误率



# Idea: Exponential Windows

- **Solution that doesn't (quite) work:**
  - Summarize **exponentially increasing** regions of the stream, looking backward
  - Drop small regions if they begin at the same point as a larger region



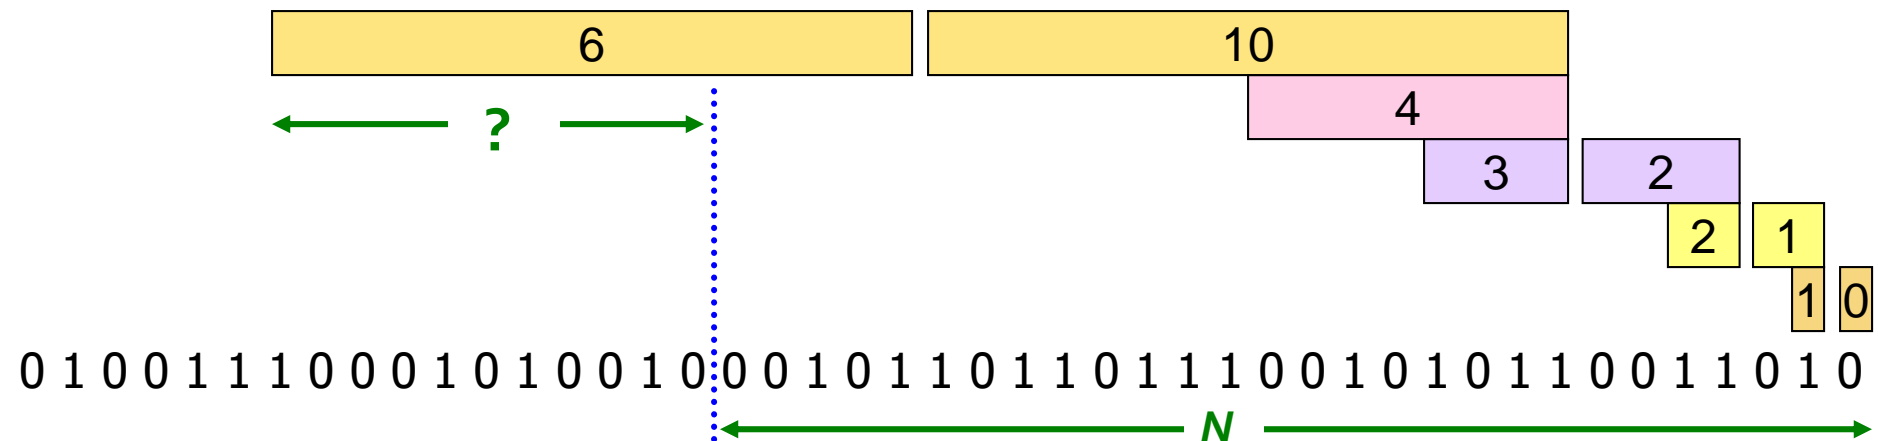
We can reconstruct the count of the last  $N$  bits, except we are not sure how many of the last 6 1s are included in the  $N$

# What's Good?

- Stores only  $O(\log^2 N)$  bits
  - $O(\log N)$  counts of  $\log_2 N$  bits each
- Easy update as more bits enter
- Error in count no greater than the number of **1s** in the “**unknown**” area


# What's Not So Good?

- As long as the **1s** are fairly evenly distributed, the error due to the unknown region is small – **no more than 50%**
- But it could be that all the **1s** are in the unknown area at the end
- In that case, **the error is unbounded!**



# Fixup: DGIM method

- **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
  - Let the block *sizes* (number of **1s**) increase exponentially
- When there are few 1s in the window, block sizes stay small, so errors are small




# DGIM: Timestamps

- Each bit in the stream has a *timestamp*, starting **1, 2, ...**
- Record timestamps modulo  **$N$**  (**the window size**), so we can represent any **relevant** timestamp in  $O(\log_2 N)$  bits

# DGIM: Buckets


- A **bucket** in the DGIM method is a record consisting of:
  - (A) The timestamp of its end [ $O(\log N)$  bits] 最近的时间
  - (B) The number of 1s between its beginning and end [ $O(\log \log N)$  bits]
- **Constraint on buckets:**  
Number of 1s must be a power of 2 只用记录指数
  - That explains the  $O(\log \log N)$  in (B) above



# Representing a Stream by Buckets

- Either **one** or **two** buckets with the same **power-of-2 number of 1s**
- Buckets do not overlap in timestamps
- Buckets are sorted by size
  - Earlier buckets are not smaller than later buckets
  - 早期bucket逐渐变大
- Buckets disappear when their end-time is  $> N$  time units in the past

# Example: Bucketized Stream



## Three properties of buckets that are maintained:

- Either **one** or **two** buckets with the same **power-of-2** number of **1s**
- Buckets do not overlap in timestamps
- Buckets are sorted by size



# Updating Buckets (1)

- When a new bit comes in, drop the last (oldest) bucket if its end-time is prior to  **$N$**  time units before the current time

**2 cases:** Current bit is **0** or **1**

- **If the current bit is 0:**  
no other changes are needed

# Updating Buckets (2)

## If the current bit is 1:

- (1) Create a new bucket of size 1, for just this bit
  - End timestamp = current time
- (2) If there are now **three buckets of size 1**,  
**combine the oldest two into a bucket of size 2**
- (3) If there are now **three buckets of size 2**,  
**combine the oldest two into a bucket of size 4**
- (4) And so on ...

# Example: Updating Buckets

Current state of the stream:

1001010110001011010101010101011010101010101110101010111010101011101010100010110010

Bit of value 1 arrives

00101011000101101010101010101101010101010111010101011101010100010110010101

Two orange buckets get merged into a yellow bucket

001010110001011010101010101011010101010101110101010111010100010110010101

Next bit 1 arrives, new orange bucket is created, then 0 comes, then 1:

0101100010110101010101010101101010101011101010101110101000101100101101

Buckets get merged...

0101100010110101010101010101101010101011101010101110101000101100101101


State of the buckets after merging

0101100010110101010101010101101010101011101010101110101000101100101101

# How to Query?


- To estimate the number of 1s in the most recent  $N$  bits:
  1. Sum the sizes of all buckets but the last  
(note “size” means the number of 1s in the bucket)
  2. Add half the size of the last bucket
- **Remember:** We do not know how many 1s of the last bucket are still within the wanted window

# Example: Bucketized Stream



# Error Bound: Proof

- **Why is error 50%? Let's prove it!**
- Suppose the last bucket has size  $2^r$
- Then by assuming  $2^{r-1}$  (i.e., half) of its **1s** are still within the window, we make an error of at most  $2^{r-1}$
- Since there is at least one bucket of each of the sizes less than  $2^r$ , the true sum is at least  $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- Thus, error at most 50%



# Further Reducing the Error


- Instead of maintaining **1** or **2** of each size bucket, we allow either  $r-1$  or  $r$  buckets ( $r > 2$ )
  - Except for the largest size buckets; we can have any number between **1** and  $r$  of those
- **Error is at most  $O(1/r)$**
- By picking  $r$  appropriately, we can tradeoff between number of bits we store and the error

# Extensions

- Can we use the same trick to answer queries

**How many 1's in the last  $k$ ?** where  $k < N$ ?

- **A:** Find earliest bucket **B** that at overlaps with  $k$ .  
Number of 1s is the **sum of sizes of more recent buckets +  $\frac{1}{2}$  size of B**




- Can we handle the case where the stream is not bits, but integers, and we want the sum of the last  $k$  elements?



# Extensions

- **Stream of positive integers**
- **We want the sum of the last  $k$  elements**
  - Amazon: Avg. price of last  $k$  sales
- **Solution:**
  - (1) If you know all have at most  $m$  bits
    - Treat  $m$  bits of each integer as a separate stream
    - Use DGIM to count 1s in each integer  $c_i$  ...estimated count for  $i$ -th bit
    - The sum is  $= \sum_{i=0}^{m-1} c_i 2^i$
  - (2) Use buckets to keep partial sums
    - Sum of elements in size  $b$  bucket is at most  $2^b$



**Idea:** Sum in each bucket is at most  $2^b$  (unless bucket has only 1 integer)  
**Bucket sizes:**

16 8 4 2 1

# Summary

- **Sampling a fixed proportion of a stream**
  - Sample size grows as the stream grows
- **Sampling a fixed-size sample**
  - Reservoir sampling
- **Counting the number of 1s in the last  $N$  elements**
  - Exponentially increasing windows
  - Extensions:
    - Number of 1s in any last  $k$  ( $k < N$ ) elements
    - Sums of integers in the last  $N$  elements